

Practical No – 01

Aim – Introduction to R tool for data analytics science.

Code –

Program 1 –

```
42+18
```

```
Arjan = 3042
```

```
class(Arjan)
```

```
print(Arjan)
```

```
print("Hello")
```

```
class(Arjan)
```

```
b=100L
```

```
class(b)
```

```
s="Afaf loves to do programming"
```

```
print(s)
```

```
class(s)
```

```
Arjan=as.integer()
```

```
class(Arjan)
```

```
x=c(42,38,52,44)
```

```
print(sqrt(x))
```

```
print(mean(x))
```

```
print(median(x))
```

```
x1=c("Arjan","Afaf","Saim","Abdul")
```

```
print(x1)
```

```
y=c(42,NA,38,NA,52,NA,44)
```

```
print(mean(y,na.rm=TRUE))
```

202001103042

Arjan Singh Johar

Btech – 1

Output –

```
> 42+18
[1] 60
> Arjan = 3042
> class(Arjan)
[1] "numeric"
> print(Arjan)
[1] 3042
> print("Hello")
[1] "Hello"
> class(Arjan)
[1] "numeric"
> b=100L
> class(b)
[1] "integer"
> s="Afaf loves to do programming"
> print(s)
[1] "Afaf loves to do programming"
> class(s)
[1] "character"
> Arjan=as.integer()
> class(Arjan)
[1] "integer"

> s="Afaf loves to do programming"
> print(s)
[1] "Afaf loves to do programming"
> class(s)
[1] "character"
> Arjan=as.integer()
> class(Arjan)
[1] "integer"
> x=c(42,38,52,44)
> print(sqrt(x))
[1] 6.480741 6.164414 7.211103 6.633250
> print(mean(x))
[1] 44
> print(median(x))
[1] 43
> x1=c("Arjan","Afaf","Saim","Abdul")
> print(x1)
[1] "Arjan" "Afaf" "Saim" "Abdul"
> y=c(42,NA,38,NA,52,NA,44)
> print(mean(y,na.rm=TRUE))
[1] 44
```

202001103042

Arjan Singh Johar

Btech – 1

Program 2 –

```
stud_data =  
data.frame(Name=c("Afaf","Arjan","Abdul"),PRN=c(38,42,44),Marks=c(100,89,89))  
  
stud_data  
  
a=20  
  
b=40  
  
if (a>b){  
    print("a is greater")  
}  
else  
    print("b is greater")  
  
x=c(1,2,3,4,5)  
  
sum=0  
  
for (i in x){  
    sum=sum+i  
}  
  
print(sum)  
  
print(sum(x))
```

Output –

```
> stud_data = data.frame(Name=c("Afaf","Arjan","Abdul"),PRN=c(38,42,44),Marks=c(100,89,89))  
> stud_data  
  Name PRN Marks  
1 Afaf  38  100  
2 Arjan 42   89  
3 Abdul 44   89  
> a=20  
> b=40  
> if (a>b){  
+   print("a is greater")  
+ }else  
+   print("b is greater")  
[1] "b is greater"  
> x=c(1,2,3,4,5)  
> sum=0  
> for (i in x){  
+   sum=sum+i  
+ }  
> print(sum)  
[1] 15  
> print(sum(x))  
[1] 15  
~
```

Practical No – 02

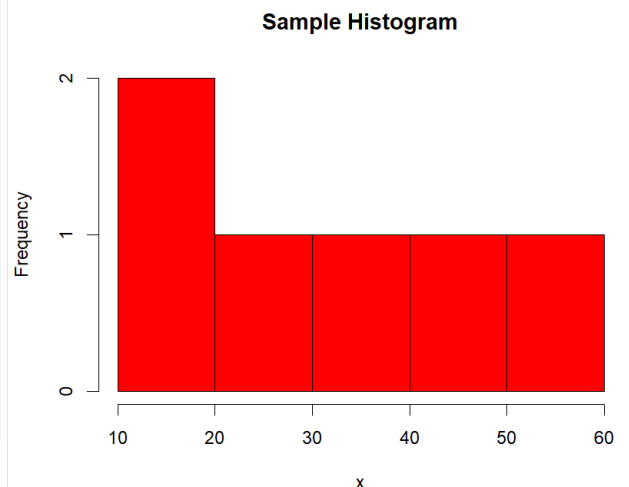

Aim – Programs for Basic Statistics and Visualization in R

Code –

Program 1 –

```
x=c(10,20,30,40,50,60)
labels=c("Arjan","Abhishek","Abdul","Afaf","Saim","Pranav")
color=c("red","green","blue","orange","purple","yellow")
pie(x,labels,col = color,main="Marks Pie Chart")
legend("topright", labels, cex = 0.8, bty = "n", fill = rainbow(length(labels)))
hist(x,col="red",main="Sample Histogram")
```

Output –



202001103042
Arjan Singh Johar
Btech – 1
Program 2 –

```
#histogram
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
samples=np.random.randint(1,100,50)
plt.hist(samples)
plt.show()
```

```
#Pie chart
slices=[70,40,50,69,35]
lang=["JAVA", "C++", "PYTHON", "C", "R"]
cols=["Red", "Green", "Blue", "Orange", "Yellow"]
plt.pie(slices, labels=lang, colors=cols, autopct="%0.2f%%", explode=[0.5,0,0,0,0])
plt.show()
```

```
#histogram
import matplotlib.pyplot as plt
import numpy as np

x = (10,20,15,14,12,25,22,11,30,35)
plt.hist(x,color="Blue")
plt.show()
```

```
import matplotlib.pyplot as plt
categories = ['Category 1', 'Category 2', 'Category 3', 'Category 4']
values = [25, 40, 30, 50]
plt.bar(categories, values)
plt.xlabel('Categories')
plt.ylabel('Values')
plt.title('Bar Plot Example')
plt.show()
```

```
import seaborn as sns
import matplotlib.pyplot as plt
sns.get_dataset_names()
tips=sns.load_dataset('tips')
tips.head()
x=[1,2,3,4,5]
plt.boxplot(x)
```

```
import seaborn as sns
```

202001103042

Arjan Singh Johar








Btech – 1

```
import matplotlib.pyplot as plt
sns.get_dataset_names()
titanic=sns.load_dataset('titanic')
titanic.head()
sns.scatterplot(x='sex',y='age',data=titanic,hue='alone')
```

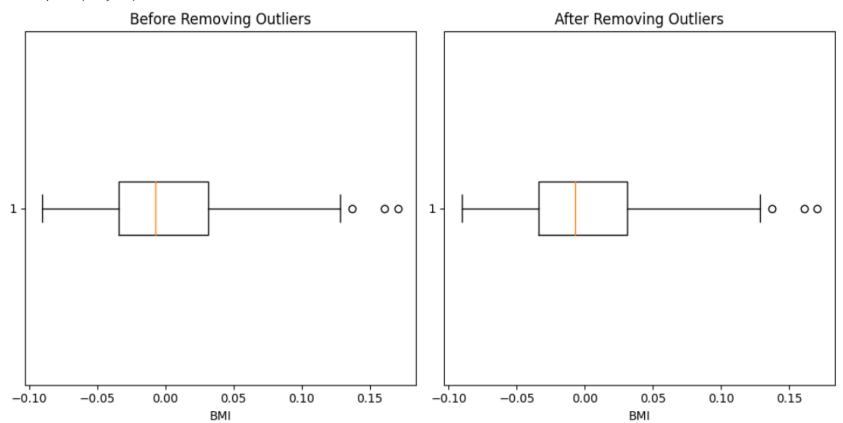
```
import matplotlib.pyplot as plt
import seaborn as sns
x=[1,2,3,4,5]
y=[10,20,30,40,50]
sns.pointplot(x=x,y=y)
```

```
import sklearn
from sklearn.datasets import load_diabetes
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
diabetes = load_diabetes()
column_name = diabetes.feature_names
df_diabetes = pd.DataFrame(diabetes.data)
df_diabetes.columns = column_name
df_diabetes.head()
print("Old Shape: ", df_diabetes.shape)
Q1 = df_diabetes['bmi'].quantile(0.25)
Q3 = df_diabetes['bmi'].quantile(0.75)
IQR = Q3 - Q1
lower = Q1 - 1.5*IQR
upper = Q3 + 1.5*IQR
upper_array = np.where(df_diabetes['bmi']>=upper)[0]
lower_array = np.where(df_diabetes['bmi']<=lower)[0]
df_diabetes.drop(index=upper_array, inplace=True)
df_diabetes.drop(index=lower_array, inplace=True)
print("New Shape: ", df_diabetes.shape)
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.boxplot(df_diabetes['bmi'], vert=False)
plt.title('Before Removing Outliers')
plt.xlabel('BMI')
plt.subplot(1, 2, 2)
plt.boxplot(diabetes.data[:, np.where(df_diabetes.columns == 'bmi')[0]], vert=False)
plt.title('After Removing Outliers')
plt.xlabel('BMI')
plt.tight_layout()
plt.show()
```

202001103042
Arjan Singh Johar
Btech – 1
Output –



Old Shape: (442, 10)
New Shape: (439, 10)



Practical No – 03

Aim – Program for simple linear regression is an approach for predicting a response using a single multiple feature.

Code –

Program 1 –


```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
data=pd.read_csv("Salary_Data.csv")
data.head()
data.info()
model = LinearRegression()
x = data[['YearsExperience']]
y = data['Salary']
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.8,random_state=0)
len(x_train)
len(x_test)
len(y_train)
len(y_test)
model.fit(x_train,y=y_train)
y_pred = model.predict(x_test)
plt.scatter(x_test,y_test,color='red',s=100)
plt.scatter(x_test,y_pred,color='blue',s=100)
plt.plot(x_test,y_pred,color='blue',linestyle="dotted")
from sklearn.metrics import r2_score
score=r2_score(y_test,y_pred)
score*100
```

Output –


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   YearsExperience  30 non-null    float64
 1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes
```


202001103042
Arjan Singh Johar
Btech – 1

[<matplotlib.lines.Line2D at 0x7e3708c2a8f0>]



<matplotlib.collections.PathCollection at 0x7e36fe7d9b70>



98.8169515729126

Program 2 –

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
data=pd.read_csv("train.csv")
data.head()
data.info()
model = LinearRegression()
x = data[['distance_traveled']]
y = data['fare']
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.8,random_state=0)
len(x_train)
len(x_test)
len(y_train)
len(y_test)
model.fit(x_train,y=y_train)
y_pred = model.predict(x_test)
y_pred
plt.scatter(x_test,y_test,color='red',s=100)
plt.scatter(x_test,y_pred,color='blue',s=100)
plt.plot(x_test,y_pred,color='blue',linestyle="dotted")
from sklearn.metrics import r2_score
score=r2_score(y_test,y_pred)
score*100
```


202001103042

Arjan Singh Johar

Btech – 1

Output –

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 209673 entries, 0 to 209672  
Data columns (total 8 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   trip_duration         209673 non-null float64  
1   distance_traveled     209673 non-null float64  
2   num_of_passengers     209673 non-null float64  
3   fare                  209673 non-null float64  
4   tip                   209673 non-null int64  
5   miscellaneous_fees    209673 non-null float64  
6   total_fare            209673 non-null float64  
7   surge_applied         209673 non-null int64  
dtypes: float64(6), int64(2)  
memory usage: 12.8 MB  
0.19829883706034002
```



Practical No – 04


Aim – Program for predicting whether a user will purchase the product or not, using logistic Regression.

Code –

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
data=pd.read_csv('Titanic-Dataset.csv')
data.head(5)
sns.countplot(x=data['Pclass'],hue=data['Survived'])
data.info()
mean_age = round((data['Age'].mean()),2)
data['Age'] = data['Age'].fillna(mean_age)
data.info()
col = ['PassengerId','Name','Cabin','Ticket','Fare']
data = data.drop(col , axis = 1)
data.head(5)
data.isnull().sum()
data = data.drop('Embarked' , axis = 1)
from sklearn.preprocessing import LabelEncoder
Encoder = LabelEncoder()
data['Sex'] = Encoder.fit_transform(data['Sex'])
data.head()
from sklearn.model_selection import train_test_split
x_train , x_test , y_train , y_test = train_test_split(x , y , train_size = 0.8 ,
random_state = 0)
len(x_train)
len(x_test)
x_test
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(x_train,y_train)
y_pred = model.predict(x_test)
y_pred
y_test
from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test , y_pred)
accuracy*100
model.predict([[1,0,31.0,1,0]])
```

202001103042
Arjan Singh Johar
Btech – 1
Output –

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   PassengerId  891 non-null    int64  
1   Survived     891 non-null    int64  
2   Pclass       891 non-null    int64  
3   Name         891 non-null    object  
4   Sex          891 non-null    object  
5   Age          714 non-null    float64  
6   SibSp        891 non-null    int64  
7   Parch        891 non-null    int64  
8   Ticket       891 non-null    object  
9   Fare         891 non-null    float64  
10  Cabin        204 non-null    object  
11  Embarked     889 non-null    object  
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB
```



```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   PassengerId  891 non-null    int64  
1   Survived     891 non-null    int64  
2   Pclass       891 non-null    int64  
3   Name         891 non-null    object  
4   Sex          891 non-null    object  
5   Age          891 non-null    float64  
6   SibSp        891 non-null    int64  
7   Parch        891 non-null    int64  
8   Ticket       891 non-null    object  
9   Fare         891 non-null    float64  
10  Cabin        204 non-null    object  
11  Embarked     889 non-null    object  
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB
```

	Survived	Pclass	Sex	Age	SibSp	Parch
0	0	3	1	22.0	1	0
1	1	1	0	38.0	1	0
2	1	3	0	26.0	0	0
3	1	1	0	35.0	1	0
4	0	3	1	35.0	0	0

	Pclass	Sex	Age	SibSp	Parch
495	3	1	29.7	0	0
648	3	1	29.7	0	0
278	3	1	7.0	4	1
31	1	0	29.7	1	0
255	3	0	29.0	0	2
...
780	3	0	13.0	0	0
837	3	1	29.7	0	0
215	1	0	31.0	1	0
833	3	1	23.0	0	0
372	3	1	19.0	0	0

179 rows × 5 columns

```
495    0  
648    0  
278    0  
31     1  
255    1  
...  
780    1  
837    0  
215    1  
833    0  
372    0  
Name: Survived, Length: 179, dtype: int64
```

array([1])

80.44692737430168

Practical No – 05

Aim – Write a program to implement k-Nearest Neighbour algorithm to classify the iris data set. Print both correct and wrong predictions

Program 1 – Titanic Dataset

Code –

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
data=pd.read_csv('Titanic-Dataset.csv')
data.head(5)
sns.countplot(x=data['Pclass'],hue=data['Survived'])
data.info()
mean_age = round((data['Age'].mean()),2)
data['Age'] = data['Age'].fillna(mean_age)
data.info()
col = ['PassengerId','Name','Cabin','Ticket','Fare']
data = data.drop(col , axis = 1)
data.head(5)
data.isnull().sum()
data = data.drop('Embarked' , axis = 1)
from sklearn.preprocessing import LabelEncoder
Encoder = LabelEncoder()
data['Sex'] = Encoder.fit_transform(data['Sex'])
data.head()
from sklearn.model_selection import train_test_split
x_train , x_test , y_train , y_test = train_test_split(x , y , train_size = 0.8 ,
random_state = 0)
len(x_train)
len(x_test)
x_test
from sklearn.neighbors import KNeighborsClassifier
model = KNeighborsClassifier(n_neighbors = 3)
model.fit(x_train , y_train)
y_pred = model.predict(x_test)
y_pred
y_test
from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test , y_pred)
accuracy*100
```

202001103042

Arjan Singh Johar

Btech – 1

Output –

```
495    0
648    0
278    0
31     1
255    1
..
780    1
837    0
215    1
833    0
372    0
Name: Survived, Length: 179, dtype: int64
```

75.41899441340783

Program 2 – Iris Dataset

Code –

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
data=pd.read_csv('Iris.csv')
data.head(5)
data.info()
data.isnull().sum()
x = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
from sklearn.model_selection import train_test_split
x_train , x_test , y_train , y_test = train_test_split(x , y ,
train_size = 0.8 , random_state = 0)
# len(x_train)
len(x_test)
from sklearn.neighbors import KNeighborsClassifier
model = KNeighborsClassifier(n_neighbors = 3)
model.fit(x_train , y_train)
y_pred = model.predict(x_test)
y_pred
y_test
from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test , y_pred)
accuracy*100
model.predict()
```

202001103042

Arjan Singh Johar

Btech – 1

Output –

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Id              150 non-null   int64
 1   SepalLengthCm   150 non-null   float64
 2   SepalWidthCm    150 non-null   float64
 3   PetalLengthCm   150 non-null   float64
 4   PetalWidthCm    150 non-null   float64
 5   Species         150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
Id              0
SepalLengthCm   0
SepalWidthCm    0
PetalLengthCm   0
PetalWidthCm    0
Species         0
dtype: int64

array(['Iris-virginica', 'Iris-versicolor', 'Iris-setosa',
      'Iris-virginica', 'Iris-setosa', 'Iris-virginica', 'Iris-setosa',
      'Iris-versicolor', 'Iris-versicolor', 'Iris-versicolor',
      'Iris-virginica', 'Iris-versicolor', 'Iris-versicolor',
      'Iris-versicolor', 'Iris-versicolor', 'Iris-setosa',
      'Iris-versicolor', 'Iris-versicolor', 'Iris-setosa', 'Iris-setosa',
      'Iris-virginica', 'Iris-versicolor', 'Iris-setosa', 'Iris-setosa',
      'Iris-virginica', 'Iris-setosa', 'Iris-setosa', 'Iris-versicolor',
      'Iris-versicolor', 'Iris-setosa'], dtype=object)
```

100.0

Practical No – 06

Aim – For the given data perform ANOVA and tell the difference of variance between groups.

Program –

Code –

```
# Load the ChickWeight dataset
data(ChickWeight)

# Subset the data to get the final weights
final_weight_data <- ChickWeight[ChickWeight$Time == max(ChickWeight$Time), ]

# Calculate the mean of final weights for each diet group
mean_weight <- tapply(final_weight_data$weight, final_weight_data$Diet, mean)

# Convert 'Diet' to a factor for ANOVA
final_weight_data$Diet <- factor(final_weight_data$Diet)

# Perform one-way ANOVA
anova_result <- aov(weight ~ Diet, data = final_weight_data)

# Post hoc test (Tukey's HSD)
tukey_result <- TukeyHSD(anova_result)

# Equality of variance (Levene's test)
library(car)

levene_test <- leveneTest(weight ~ Diet, data = final_weight_data)

print(levene_test)

# Create a boxplot
boxplot(weight ~ Diet, data = final_weight_data, main = "Boxplot of Final Weight by Diet", xlab = "Diet", ylab = "Final Weight")

# Check the p-value from ANOVA
p_value <- summary(anova_result)[[1]][["Pr(>F)"]][1]

# Define the significance level (alpha)
alpha <- 0.05
```



```
# Print the decision based on the p-value
```

```
if (p_value < alpha) {
```


```
  cat("Reject the null hypothesis. There is a significant difference in means.\n")
```

```
} else {
```

```
  cat("Accept the null hypothesis. There is no significant difference in means.\n")
```

```
}
```

Output –



```
Levene's Test for Homogeneity of Variance (center = median)
```

```
  Df F value Pr(>F)
```

```
group 3  1.1597 0.3367
```

```
  41
```

```
Reject the null hypothesis. There is a significant difference in means.
```

Practical No – 07

Aim – Implementation of Naive Bayesian Classifier.

Code and Output –

```
library(e1071)

# Load the Iris dataset

data(iris)

library(caret)

set.seed(123) # For reproducibility

train_index <- createDataPartition(iris$Species, p = 0.7, list = FALSE)

train_data <- iris[train_index, ]

test_data <- iris[-train_index, ]

# Train the Naive Bayes classifier

nb_classifier <- naiveBayes(Species ~ ., data = train_data)

# Predict on the test set

test_predictions <- predict(nb_classifier, newdata = test_data)

# Create a confusion matrix

confusion_matrix <- table(Actual = test_data$Species, Predicted = test_predictions)

confusion_matrix

# Calculate accuracy

accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)

accuracy
```

Output –

```
-----
Actual      Predicted
setosa      setosa versicolor virginica
setosa      15         0         0
versicolor  0         13         2
virginica   0         2         13
>
> # Calculate accuracy
> accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
> accuracy
[1] 0.9111111
> |
```

Practical No – 08

Aim – Program for market basket analysis using Association Rules.

Program –

```
# Load the libraries
install.packages("arulesViz")
library(arules)
library(arulesViz)
library(datasets)

# Load the data set
data(Groceries)

# Create an item frequency plot for the top 20 items
itemFrequencyPlot(Groceries,topN=20,type="absolute")

# Get the rules
rules <- apriori(Groceries, parameter = list(supp = 0.001, conf = 0.8))

# Show the top 5 rules, but only 2 digits
options(digits=2)
inspect(rules[1:5])

rules<-sort(rules, by="confidence", decreasing=TRUE)

rules <- apriori(Groceries, parameter = list(supp = 0.001, conf = 0.8,maxlen=3))
subset.matrix <- is.subset(rules, rules)
subset.matrix[lower.tri(subset.matrix, diag=T)] <- NA
redundant <- colSums(subset.matrix, na.rm=T) >= 1
rules.pruned <- rules[!redundant]
rules<-rules.pruned

rules<-apriori(data=Groceries, parameter=list(supp=0.001,conf = 0.08),
              appearance = list(default="lhs",rhs="whole milk"),
              control = list(verbose=F))
```

202001103042

Arjan Singh Johar

Btech – 1

```
rules<-sort(rules, decreasing=TRUE,by="confidence")
```

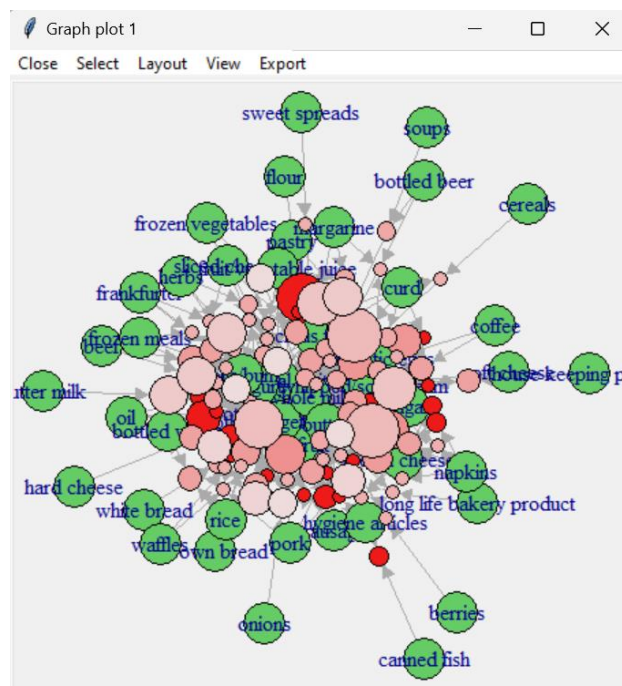

```
inspect(rules[1:5])
```

```
library(arulesViz)
```

```
plot(rules,method="graph",interactive=TRUE)
```

Output –

```
lhs              rhs      support confidence coverage lift
[1] {rice, sugar}   => {whole milk} 0.0012 1      0.0012 3.9
[2] {canned fish, hygiene articles} => {whole milk} 0.0011 1      0.0011 3.9
[3] {root vegetables, butter, rice} => {whole milk} 0.0010 1      0.0010 3.9
[4] {root vegetables, whipped/sour cream, flour} => {whole milk} 0.0017 1      0.0017 3.9
[5] {butter, soft cheese, domestic eggs} => {whole milk} 0.0010 1      0.0010 3.9
count
[1] 12
[2] 11
[3] 10
[4] 17
[5] 10
```




Practical No – 09

Aim – Program to implement a Decision tree algorithm for any application.

Code –

```
# Load the required libraries  
library(rpart)  
library(rpart.plot)  
# Load the "mtcars" dataset (built-in dataset in R)  
data(mtcars)  
# Split the dataset into features (X) and target (y)  
X <- mtcars[, -1] # Features (excluding the first column, which is the car model)  
y <- mtcars$mpg # Target (miles per gallon)  
# Build the decision tree model  
model <- rpart(y ~ ., data = data.frame(X, y), method = "anova")  
# Visualize the decision tree  
rpart.plot(model, box.palette = "auto", type = 3, extra = 1)
```

Output –



Practical No – 10

Aim – Program to Simulate Principal component analysis.

Code –

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
# Load the Iris dataset
iris = load_iris()
X = iris.data
y = iris.target
# Standardize the features (mean=0 and variance=1)
X_standardized = StandardScaler().fit_transform(X)
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_standardized, y,
test_size=0.4, random_state=42)
# Create a Naive Bayes classifier
clf = GaussianNB()
# Train the classifier
clf.fit(X_train, y_train)
# Make predictions on the test data
predictions = clf.predict(X_test)
from sklearn.metrics import accuracy_score
accuracy_score(y_test, predictions)
pca = PCA(n_components=1)
X_pca = pca.fit_transform(X_standardized)
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_pca, y,
test_size=0.4, random_state=42)
# Create a Naive Bayes classifier
clf = GaussianNB()
# Train the classifier
clf.fit(X_train, y_train)
# Make predictions on the test data
predictions = clf.predict(X_test)
from sklearn.metrics import accuracy_score
accuracy_score(y_test, predictions)
```

202001103042


Arjan Singh Johar

Btech – 1

```
# Apply PCA and reduce the data to 2 principal components
pca = PCA(n_components=)
X_pca = pca.fit_transform(X_standardized)
# Print the explained variance ratio
print("Explained Variance Ratio (2 components):",
pca.explained_variance_ratio_)
# Plot the data points in the reduced feature space
plt.figure(figsize=(8, 6))
colors = ['blue', 'red', 'green']
for i, target_name in enumerate(iris.target_names):
    plt.scatter(X_pca[y == i, 0], X_pca[y == i, 1], color=colors[i],
label=target_name)
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('PCA on Iris Dataset')
plt.legend()
plt.show()
```

Output –

0.9666666666666667



Practical No – 11

Aim – Program to Simulate Principal component analysis

Code –

```
# Load the USArrests dataset
data("USArrests")

# Perform SVD on the dataset
svd_result <- svd(USArrests)

# Extract U, Sigma, and V matrices
U <- svd_result$u
Sigma <- diag(svd_result$d)
V <- svd_result$v

# Print the matrices
print("U (left singular vectors):")
print(U)
print("Sigma (singular values):")
print(Sigma)
print("V (right singular vectors):")
print(V)
```

Output –

```
> print("U (left singular vectors):")
[1] "U (left singular vectors):"
> print(U)
      [,1]      [,2]      [,3]      [,4]
[1,] -0.17162510  0.096325710  0.065154797  0.153695511
[2,] -0.18911657  0.173452566 -0.426657848 -0.178014378
[3,] -0.21559302  0.078998111  0.020637399 -0.280707843
[4,] -0.13902443  0.059889811  0.013922695  0.016104178
[5,] -0.20677884 -0.009812026 -0.176332443 -0.218674250
[6,] -0.15587942 -0.064555293 -0.282882796 -0.117974190
[7,] -0.09086363 -0.196817368  0.177814176 -0.056150268
[8,] -0.17536307  0.035102548  0.242423936 -0.223770615
[9,] -0.24315375  0.146502368  0.050477542  0.025718639
[10,] -0.15591071  0.042885364 -0.069631843  0.426192214
[11,] -0.05035785 -0.336841681 -0.093988180  0.169255594
[12,] -0.09273525 -0.071651205  0.048571905 -0.144733647
[13,] -0.18583902 -0.004760115  0.112681109 -0.023621705
[14,] -0.09113246 -0.140219345 -0.077396606  0.106957520
[15,] -0.05057860 -0.189585706  0.028511452 -0.008876337
[16,] -0.09241257 -0.139884238 -0.004741157  0.048135122
[17,] -0.08535772 -0.080906191 -0.029723458  0.262636519
[18,] -0.18215443  0.078717908  0.086540399  0.247322269
[19,] -0.06696147 -0.113964054  0.123029673 -0.066578837
[20,] -0.21660706  0.153849690  0.049568029 -0.114685718
```

```
>
> print("Sigma (singular values):")
[1] "Sigma (singular values):"
> print(Sigma)
      [,1]      [,2]      [,3]      [,4]
[1,] 1419.061  0.0000  0.00000  0.00000
[2,]  0.000 194.8258  0.00000  0.00000
[3,]  0.000  0.0000 45.66134  0.00000
[4,]  0.000  0.0000  0.00000 18.06956
>
> print("V (right singular vectors):")
[1] "V (right singular vectors):"
> print(V)
      [,1]      [,2]      [,3]      [,4]
[1,] -0.04239181  0.01616262 -0.06588426  0.99679535
[2,] -0.94395706  0.32068580  0.06655170 -0.04094568
[3,] -0.30842767 -0.93845891  0.15496743  0.01234261
[4,] -0.10963744 -0.12725666 -0.98347101 -0.06760284
```


Practical No – 12

Aim – Using built-in function perform Support Vector machine algorithm

Code –

```
library(e1071)

data(iris)

set.seed(123) # for reproducibility

sample_indices <- sample(1:nrow(iris), nrow(iris) * 0.7) # 70% for training

train_data <- iris[sample_indices, ]

test_data <- iris[-sample_indices, ]

svm_model <- svm(Species ~ ., data = train_data, kernel = "linear")

predictions <- predict(svm_model, test_data)

confusion_matrix <- table(predictions, test_data$Species)

print(confusion_matrix)

print(predictions)

accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)

cat("Accuracy:", accuracy, "\n")
```

Output –

```
predictions  setosa versicolor virginica
setosa      14         0         0
versicolor   0        17         0
virginica    0         1        13
> print(predictions)
 1  2  3  5 11 18 19
setosa setosa setosa setosa setosa setosa setosa
28 29 33 36 45 48 49
setosa setosa setosa setosa setosa setosa setosa
55 56 57 58 59 61 62
versicolor versicolor versicolor versicolor versicolor versicolor versicolor
65 66 68 70 77 83 84
versicolor versicolor versicolor versicolor versicolor versicolor virginica
94 95 98 100 101 104 105
versicolor versicolor versicolor versicolor virginica virginica virginica
111 113 116 125 131 133 135
virginica virginica virginica virginica virginica virginica virginica
140 141 145
virginica virginica virginica
Levels: setosa versicolor virginica
> accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
> cat("Accuracy:", accuracy, "\n")
Accuracy: 0.9777778
```